


Lecture 7

Binary Search Trees and Red-Black Trees

Motivation for binary search trees

- We've been assuming that we have access to some basic data structures:
 - (Sorted) linked lists




- (Sorted) arrays




Sorted linked lists



- $O(1)$ insert/delete (assuming we have a pointer to the location of the insert/delete):



- $O(n)$ search/select:




Sorted Arrays



- $O(n)$ insert/delete:



- $O(\log(n))$ search, $O(1)$ select:









Search: Binary search to see if 3 is in A.

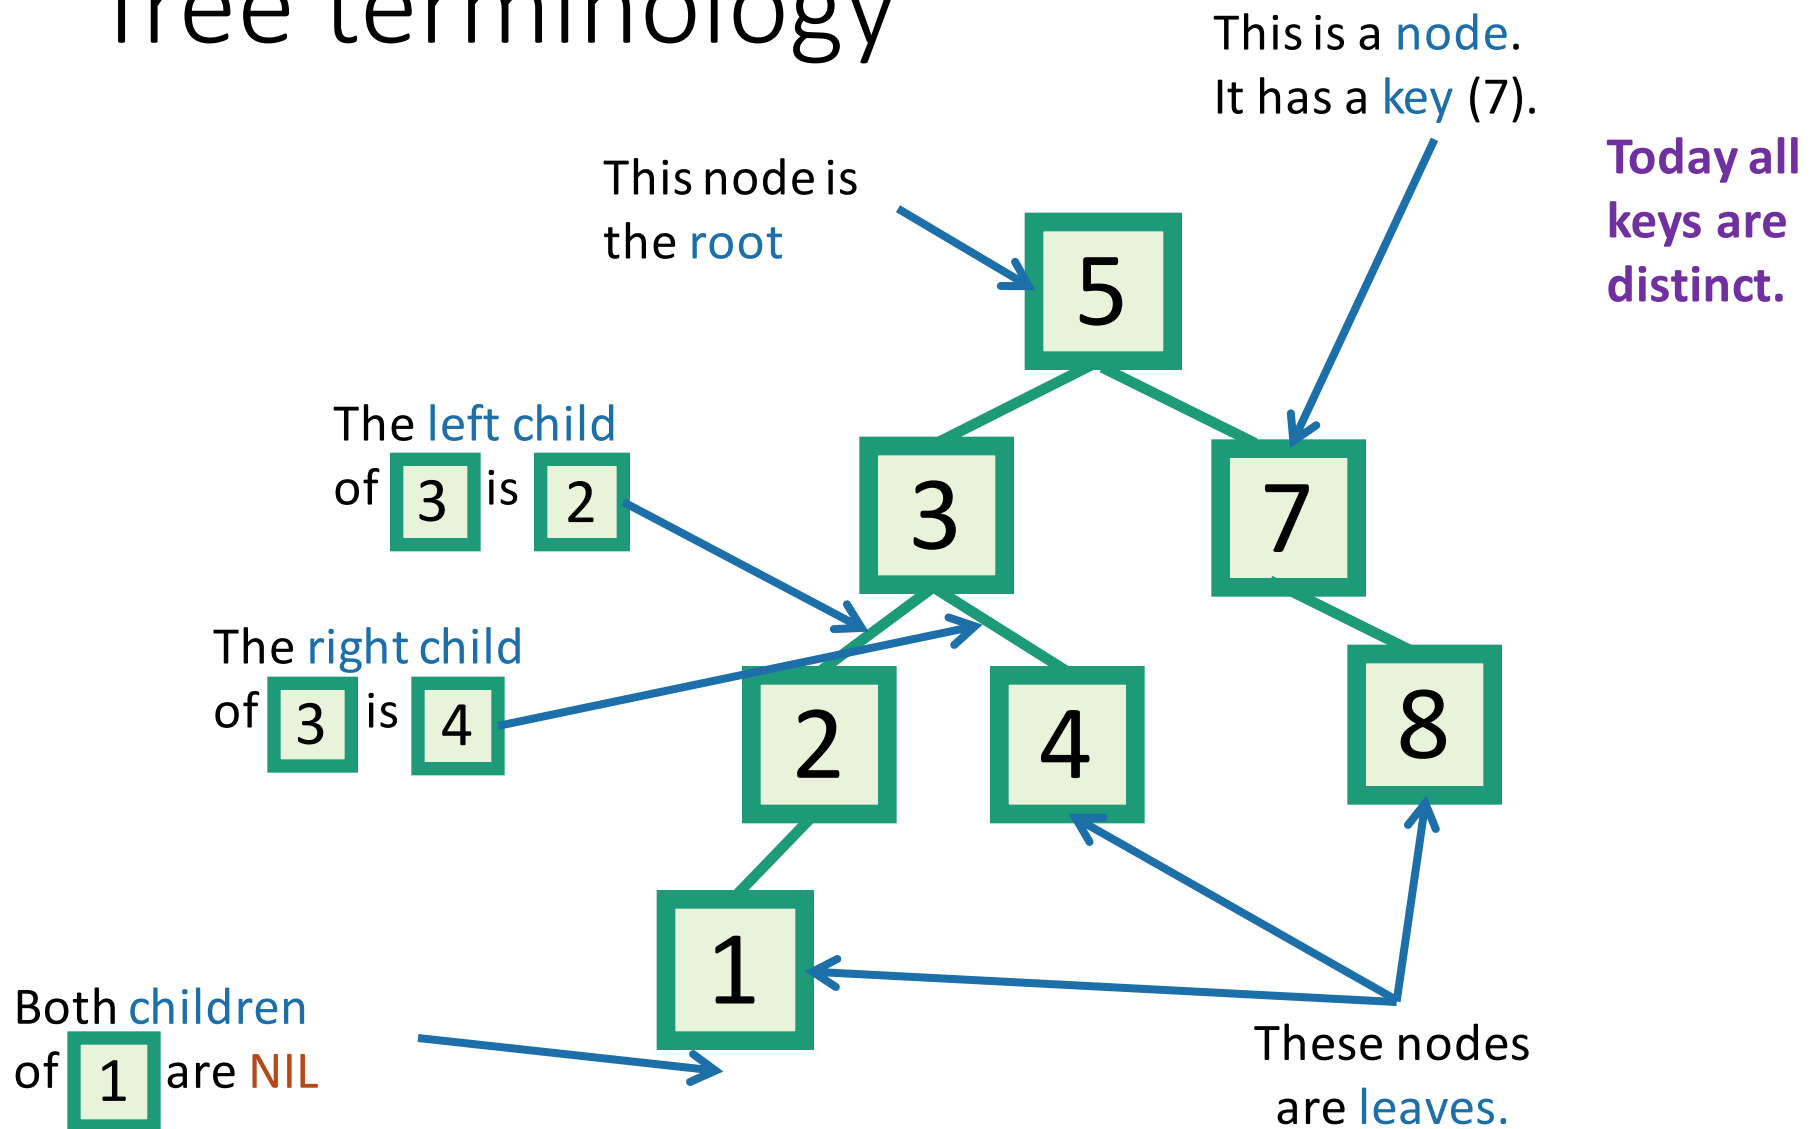
Select: Third smallest is A[3].

The best of both worlds

TODAY!


	Sorted Arrays	Linked Lists	Binary Search Trees*
Search	$O(\log(n))$ 	$O(n)$ 	$O(\log(n))$ 
Insert/Delete	$O(n)$ 	$O(1)$ 	$O(\log(n))$ 

Tree terminology




Binary Search Trees

- It's a **binary tree** so that:
 - Every LEFT descendent of a node has key less than that node.
 - Every RIGHT descendent of a node has key larger than that node.
- Example of building a binary search tree:




Binary Search Trees

- It's a **binary tree** so that:
 - Every LEFT descendent of a node has key less than that node.
 - Every RIGHT descendent of a node has key larger than that node.
- Example of building a binary search tree:




Binary Search Trees

- It's a **binary tree** so that:
 - Every LEFT descendent of a node has key less than that node.
 - Every RIGHT descendent of a node has key larger than that node.
- Example of building a binary search tree:




Binary Search Trees

- It's a **binary tree** so that:
 - Every LEFT descendent of a node has key less than that node.
 - Every RIGHT descendent of a node has key larger than that node.
- Example of building a binary search tree:




Binary Search Trees

- It's a **binary tree** so that:
 - Every LEFT descendent of a node has key less than that node.
 - Every RIGHT descendent of a node has key larger than that node.
- Example of building a binary search tree:





Aside: this should look familiar
kinda like QuickSort




Binary Search Trees

- It's a **binary tree** so that:
 - Every **LEFT** descendent of a node has key less than that node.
 - Every **RIGHT** descendent of a node has key larger than that node.



SEARCH in a Binary Search Tree

definition by example




EXAMPLE: Search for 4.

EXAMPLE: Search for 4.5

- It turns out it will be convenient to **return 4** in this case
- (that is, **return** the last node before we went off the tree)


INSERT in a Binary Search Tree



EXAMPLE: Insert 4.5

- **INSERT(key):**
 - $x = \text{SEARCH}(\text{key})$
 - **if** $\text{key} > x.\text{key}$:
 - Make a new node with the correct key, and put it as the right child of x .
 - **if** $\text{key} < x.\text{key}$:
 - Make a new node with the correct key, and put it as the left child of x .
 - **if** $x.\text{key} == \text{key}$:
 - **return**

DELETE in a Binary Search Tree




EXAMPLE: Delete 2

- **DELETE**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **if** $x.\text{key} == \text{key}$:
 -delete x


← This is a bit more complicated...

How long do these operations take?

- **SEARCH** is the big one.
- Everything else just calls **SEARCH** and then does some small $O(1)$ -time operation.



Wait...



- This is a valid binary search tree.
- The version with n nodes has depth n , **not** $O(\log(n))$.

Could such a tree show up?
In what order would I have to
insert the nodes?

Inserting in the order
2,3,4,5,6,7,8 would do it.

So this could happen.




What to do?

- Keep track of how deep the tree is getting.



Idea 1: Rotations

No matter what lives underneath A,B,C,
this takes time $O(1)$. (Why?)

- Maintain Binary Search Tree (BST) property, while moving stuff around.



This seems helpful



Idea 2: have some proxy for balance

- Maintaining perfect balance is too hard.
- Instead, come up with some proxy for balance:
 - If the tree satisfies [SOME PROPERTY], then it's pretty balanced.
 - We can maintain [SOME PROPERTY] using rotations.

Red-Black Trees

- A Binary Search Tree that balances itself!
- No more time-consuming by-hand balancing!
- Be the envy of your friends and neighbors with the time-saving...

Red-Black tree!

Maintain balance by stipulating that **black nodes** are balanced, and that there aren't too many **red nodes**.


It's just good sense!



Red-Black Trees





these rules are the proxy for balance

- Every node is colored **red** or **black**.
- The root node is a **black node**.
- NIL children count as **black nodes**.
- Children of a **red node** are **black nodes**.
- For all nodes x:
 - all paths from x to NONE's have the same number of black nodes on them.



Examples(?)


- Every node is colored **red** or **black**.
- The root node is a **black node**.
- **NONE** children count as **black nodes**.
- Children of a **red node** are **black nodes**.
- For all nodes x:
 - all paths from x to NONE's have the same number of black nodes on them.



That turns out to be basically right.

[proof sketch]

- Say there are $b(x)$ black nodes in any path from x to $NONE$.
 - (including x).
- Claim:
 - Then there are at least $2^{b(x)} - 1$ nodes in the subtree underneath x .
- [Proof by induction — on board if time]



Then:

$$\begin{aligned} n &\geq 2^{b(\text{root})} - 1 && \text{using the Claim} \\ &\geq 2^{\text{height}/2} - 1 && b(\text{root}) \geq \text{height}/2 \text{ because of RBTree rules.} \end{aligned}$$

Rearranging:

$$n + 1 \geq 2^{\text{height}/2} \Rightarrow \text{height} \leq 2\log(n + 1)$$


Okay, so it's balanced...

...but can we maintain it?


- Yes!

Inserting into a Red-Black Tree

- Make a new **red node**.
- Insert it as you would normally.



What if it looks like this?




Example: insert 0




Inserting into a Red-Black Tree

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.




What if it looks like this?

Example: insert 0




Inserting into a Red-Black Tree

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.



What if it looks like this?



Example: insert 0



Can't we just insert 0 as a **black node**?



Inserting into a Red-Black Tree

-1

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.




What if it looks like this?

Example: insert 0


- Instead recolor like this.
- Need to argue:
 - RB-Tree properties still hold.
- What about the red root?
 - if 6 is actually the root, color it black.
 - Else, recursively re-color up the tree.



Now the problem looks like this, where I'm inserting




Inserting into a Red-Black Tree



What if it looks like this?

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.






Example: Insert 0.

- Actually, **this can't happen?**
- It might happen that we just turned 0 red from the previous step.
- Or it could happen if **7** is actually NIL.


Recall Rotations

- Maintain Binary Search Tree (BST) property, while moving stuff around.





Inserting into a Red-Black Tree

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.



What if it looks like this?

YOINK!



Need to argue that if RB-Tree property held before, it still does.







That's basically it

- A few things still left to check for **INSERT!**
 - Anything else that might happen looks basically like what we just did.
 - Formally dealing with the recursion.
 - **You check these!** (or see CLRS)
- **DELETE** is similar.

The punchline:

- Red-Black Trees **always** have height at most $2\log(n+1)$.
- As with general **Binary Search Trees**, all operations are $O(\text{height})$
- **So all operations are $O(\log(n))$.**

Conclusion: The best of both worlds

	Sorted Arrays	Linked Lists	Balanced Binary Search Trees
Search	$O(\log(n))$ 	$O(n)$ 	$O(\log(n))$ 
Insert/Delete	$O(n)$ 	$O(1)$ 	$O(\log(n))$ 

Summary

- **Balanced binary trees** are the best of both worlds!
- But we need to **keep them balanced**.
- **Red-Black Trees** do that for us.
 - We get $O(\log(n))$ -time INSERT/DELETE/SEARCH
 - Clever idea: have a **proxy for balance**