

Multi-probe consistent hashing

Ben Appleton, Michael O'Reilly
{appleton, moreil}@google.com
Google

May 4, 2015

1 Abstract

We describe a consistent hashing algorithm which performs multiple lookups per key in a hash table of nodes. It requires no additional storage beyond the hash table, and achieves a peak-to-average load ratio of $1 + \varepsilon$ with just $1 + \frac{1}{\varepsilon}$ lookups per key.

2 Introduction

Consistent hashing was introduced by Karger et al. [4]. It allows multiple clients to balance keys over a set of nodes without communication, and continue to agree for almost every key as the collection of live nodes changes over time or varies across machines. First applied to caching, consistent hashing has also been applied to key-value stores such as Dynamo [2], and routing tables such as Kademlia [6], used notably by BitTorrent.

A figure of merit for load balancing is the *peak-to-average load ratio*. In this paper we consider the case where there are many distinct keys per node, and define the load on a node as the proportion of keys which map to a node. Then peak-to-average load ratio is the ratio of the maximum load to the average load over the set of nodes.

In addition to consistent placement, an ideal consistent hash might have the following performance properties¹:

1. $O(n)$ space for n nodes; ideally just the collection of nodes.
2. $O(1)$ time per insertion or removal.
3. $O(1)$ time per lookup.
4. Peak-to-average load ratio at most $1 + \varepsilon$, for some small ε .

However existing algorithms fall short of this ideal.

Karger et al's *ring consistent hash* [4] hashes each node $O(\frac{\ln n}{\varepsilon^2})$ ways to a ring, indexing each node hash. To assign a key to a node, it hashes the key to the ring and returns the node with the next hash. However to obtain a peak-to-average load ratio of $1 + \varepsilon$ it requires $O(\frac{n \ln n}{\varepsilon^2})$ memory.

Thaler and Ravishankar's highest random weight algorithm [10] hashes each key against each node, returning the node with the highest resulting hash. For a large number of keys this produces a peak-to-average load ratio of 1, but takes $O(n)$ time per lookup. Wang and Ravishankar [11] present a variation which takes $O(\ln n)$ time, by clustering nodes into a pre-agreed tree then recursively selecting clusters by hashing the key against each cluster down the tree. However this requires pre-agreement on the hierarchy, with no provision for handling changes to the set of live nodes.

Lamping and Veach's *jump consistent hash* [5] hashes each key to a list of nodes labeled $1, 2, \dots, n$. Keys are placed using a pseudo-random number generator to compute a sequence of node assignments as the number of nodes grows. It takes $O(1)$ space and $O(\ln n)$ expected time, achieving a peak-to-average load ratio of exactly 1. However, it does not

¹Here $O(\dots)$ should be interpreted liberally to allow in expectation, with high probability, amortized, or *etc.*

support the removal of arbitrary nodes. This prevents from using jump consistent hash in applications which must handle arbitrary node loss. It also prevents from using jump consistent hash in weighted consistent hashing [9].

3 Analysis

In this paper we define ‘with high probability’ to be with probability $1 - \frac{1}{n^{\Omega(1)}}$.

3.1 Hashing keys and nodes once each - naive approach

Consider hashing n nodes to the unit ring. When presented with a key, hash the key to the unit ring, and return the next node along the ring. This requires $O(n)$ memory and $O(1)$ time per lookup, but produces a poor peak-to-average load.

Let X_1, \dots, X_n be the distances between successive node hashes, such that the probability of selecting node i is X_i . For a node hash function selected at random from a universal ranged hash family, the distances are distributed with $\Pr(X_i > x_i) = (1 - x_i)^n$. As Feller shows ([3] Chapter I.7), in the limit $n \rightarrow \infty$ the distances converge to a Poisson process with independent and identically distributed (iid) $X_i \sim \text{Exp}(n)$; a simplifying approximation which we use subsequently. Therefore with high probability $\max_{i=1}^n X_i = \Theta(\frac{\ln n}{n})$. Since the mean load is $\frac{1}{n}$, the peak-to-average load ratio is then $\Theta(\ln n)$. Since a service must be provisioned for peak load, but its capacity is proportional to the average load, a high peak-to-average load ratio may be unacceptable in many applications.

3.2 Hashing nodes J ways

Ring consistent hash [4] resolves this load imbalance by hashing each node $J = O(\ln n)$ ways to *virtual nodes* on the unit ring. The virtual nodes are stored in a hash table. When presented with a key, hash the key to the unit ring, find the next virtual node, then return the corresponding physical node.

If we use J independent hashes, the set of node hashes forms a Poisson process with iid $X_{i,j} \sim \text{Exp}(Jn)$, and fraction of keys assigned to node j is $S_j = \sum_{i=1}^J X_{i,j}$ with mean $\frac{1}{n}$. By Cramer’s theorem [1] $P(S_j > \frac{1+\varepsilon}{n}) \sim e^{-JI(1+\varepsilon)}$ where $I(t) = Jnt - 1 - \ln(Jnt)$ for this process. So to achieve a peak-to-average load ratio of $1 + \varepsilon$ in expectation or with high probability requires $J = \Theta(\frac{\ln n}{\varepsilon^2})$.

Note that J cannot be changed online as that would break consistency. So J must be sized for the maximum number of nodes expected in the lifetime of the system, or provision must be made for changes that break consistency.

3.3 Hashing keys K ways

We propose to store each node once but to hash keys K ways, returning the subsequent node which is closest to a key hash. We call this *multi-probe consistent hashing*. This requires $O(n)$ space to store n nodes and $O(K)$ time per lookup. Perhaps surprisingly, using only $K = 2$ key hashes improves the peak-to-average load ratio from $O(\ln n)$ to $O(1)$ with high probability.

The key result is Theorem 1. For technical reasons we begin by deriving the expected peak-to-average load ratio.

Lemma 1. *For K independent hashes per key with $2 \leq K \ll \frac{\sqrt{n}}{\ln n}$, for a random node hash function from a universal ranged hash family, the peak load is $\frac{K}{K-1} \frac{1}{n} + o\left(\frac{1}{n}\right)$ in expectation.*

Proof. Consider a $K + 1$ -independent universal ranged hash family, selecting 1 node hash function and K key hash functions. Then the node hashes form a Poisson process with rate n . Without loss of generality let $x_1 = \max_{i=1}^n x_i$, such that node 1 is the maximally loaded node. For $K \ll \frac{\sqrt{n}}{\ln n}$ the probability that multiple key hashes resolve to x_1 is $o\left(\frac{1}{n}\right)$, which case we will neglect. Else if 1 key hash resolves to x_1 it has distance $\sim U(0, x_1)$ and the $K - 1$ other key hashes have iid distances $\sim \text{Exp}(n + 1)$, where the latter is obtained by considering the key hash as another node hash. Then the probability that a key is assigned to x_1 is:

$$\begin{aligned} & K \int_{x=0}^1 e^{-(n+1)(K-1)x} dx + o\left(\frac{1}{n}\right) \\ &= \frac{K}{K-1} \frac{1}{n} + o\left(\frac{1}{n}\right) \end{aligned} \tag{1}$$

□

In [7] McDiarmid proved:

Lemma 2. *Let X_1, \dots, X_n be a family of independent random variables. Suppose that the real-valued function Z satisfies*

$$|Z(x) - Z(x')| \leq c_k \tag{2}$$

whenever the vectors x and x' differ only in the k th coordinate. Let μ be the expected value of the random variable $Z(X)$. Then for any $\lambda \geq 0$,

$$\Pr(|Z(X) - \mu| \geq \lambda \sigma) \leq 2e^{-2\lambda^2} \tag{3}$$

where $\sigma^2 = \sum_{k=1}^n c_k^2$.

We proceed to use McDiarmid's inequality to prove that the bound in Lemma 1 holds not just in expectation but with high probability.

Theorem 1. For K independent hashes per key with $2 \leq K \ll \frac{\sqrt{n}}{\ln n}$, for a random node hash function from a universal ranged hash family, the peak load is $\frac{K}{K-1} \frac{1}{n} + o\left(\frac{1}{n}\right)$ with high probability.

Proof. For one independent key hash, the probability that the distance to the next node hash is at most x is $F(x)$, with:

$$1 - F(x) = \sum_{i=1}^n \begin{cases} x_i - x & \text{if } x \leq x_i \\ 0 & \text{if } x > x_i \end{cases} \quad (4)$$

For $K \ll \frac{n}{\ln n}$ the probability that a single key hashes multiple times to the maximally loaded node is $o\left(\frac{1}{n}\right)$. Then defining

$$Z_K = K \int_{x=0}^1 (1 - F(x))^{K-1} dx \quad (5)$$

the peak load for K key hashes is $Z_K + o\left(\frac{1}{n}\right)$. For a random node hash function, the expected value of Z_K is $\frac{K}{K-1} \frac{1}{n} + o\left(\frac{1}{n}\right)$.

Recall we have $\mu = \frac{K}{K-1} \frac{1}{n} + o\left(\frac{1}{n}\right)$, and $0 \leq x_k \leq \frac{c \ln n}{n}$ with high probability.

Begin with the case $K = 2$. Equation 5 simplifies to $Z_2 = \sum_{i=1}^n x_i^2$, which gives $c_k = \left(\frac{c \ln n}{n}\right)^2$ and hence $\sigma = \frac{(c \ln n)^2}{n \sqrt{n}} = o\left(\frac{1}{n}\right)$, so $Z_2 = \frac{2}{n} + o\left(\frac{1}{n}\right)$ with high probability.

For the case $K > 2$ we obtain

$$\begin{aligned} c_k &= K \frac{c \ln n}{n} (K-1) \int_{x=0}^1 (1 - F(x))^{K-2} dx \\ &= K \frac{c \ln n}{n} Z_{K-1} \\ &= O\left(\frac{K \ln n}{n^2}\right) \end{aligned} \quad (6)$$

The first equality notes that x_k appears $K-1$ times in $(1 - F(x))^{K-1}$, whose difference with respect to x_k is at most $\frac{c \ln n}{n} (K-1) (1 - F(x))^{K-2}$, discarding higher-order terms by $K = o\left(\frac{n}{\ln n}\right)$. We then induct on K . Then by $\sigma^2 = \sum_{k=1}^n c_k^2$ we have $\sigma = O\left(\frac{K \ln n}{n \sqrt{n}}\right)$. Since $K = o\left(\frac{\sqrt{n}}{\ln n}\right)$ we obtain $\sigma = o\left(\frac{1}{n}\right)$, and hence the desired result. \square

So the peak-to-average load ratio is $\frac{K}{K-1} + o(1)$ with high probability. To achieve a peak-to-average load ratio of $1 + \varepsilon$ requires $K = 1 + \frac{1}{\varepsilon}$ hashes, and $O\left(\frac{1}{\varepsilon}\right)$ time per lookup.

Table 1: Properties of each algorithm

	Jump c. h.	Ring c. h.	Multi-probe c. h.
Peak-to-average	1	$1 + \varepsilon$	$1 + \varepsilon$
Memory	$O(1)$	$O(\frac{n \ln n}{\varepsilon^2})$	$O(n)$
Update time	0	$O(\frac{\ln n}{\varepsilon^2})$	$O(1)$
Assignment time	$O(\ln n)$	$O(1)$	$O(\frac{1}{\varepsilon})$
Arbitrary node removal	No	Yes	Yes

3.4 Other properties

Karger et al [4] defined other important properties for consistent hash functions: *monotonicity*, *spread* and *load*. For completeness we will address these properties. Our proofs closely mirror [4].

Theorem 2. *For $K = O(1)$ the hash family F described in this paper has the following properties:*

1. *Monotonicity: F is monotone.*
2. *Spread: If the number of views $V = \rho n$ for some constant ρ , and the number of keys $I = n$, then for $i \in \mathcal{I}$, $\sigma(i)$ is $O(t \ln n)$ with high probability.*
3. *Load: If V and I are as above, then for $n \in \mathcal{N}$, $\lambda(n)$ is $O(t \ln n)$ with high probability.*

Proof. Monotonicity: Adding a node to the ring does not increase the distance from any key hash to the next node, and does not reduce the distance from any key hash to any existing node. So no key can switch to an existing node.

Spread and load follow from the observation that with high probability, a point from every view falls into an interval of length $O(\frac{t \ln n}{n})$. Spread follows by observing that for each key, the number of node points that fall within this size interval around the K key hashes, $O(t \ln n)$, is an upper bound on the spread of that key. Load follows by counting the number of key hashes that fall in the region owned by a node hash, $O(t \ln n)$. \square

3.5 Performance summary

Table 1 summarizes the performance of each algorithm for n nodes.

Table 2: Peak-to-average, $K = 2$

Number of nodes	median of trials	90%ile of trials	99%ile of trials
10	1.74	2.43	3.32
100	1.96	2.22	2.48
1,000	2.00	2.08	2.16
10,000	2.00	2.03	2.05
100,000	2.00	2.01	2.02

4 Implementation

For a hash table we use an array of sorted, inlined vectors. The array is sized to about 6 nodes per inlined vector. The inlined vector stores the first 8 elements inline, then spills to an out-of-line buffer. This avoids pointer-chasing in the common case. We use 64-bit identifiers and hashes. We store the hash alongside the node identifier to save on subsequent hashes.

5 Performance

We compare multi-probe consistent hash to ring consistent hash [4] and jump consistent hash [5].

5.1 Peak-to-average load ratio

We measured the peak-to-average load ratio over a range of node counts. For each node count we ran 1,000 trials using different node hash seeds to obtain percentiles over the statistic of interest: peak-to-average load ratio. For each trial we sampled 1,000,000 keys per node. These simulations were run on a cluster of machines.

Table 2 shows the peak-to-average load ratio for multi-probe consistent hash with $K = 2$. The peak-to-average load ratio converges to 2. This requires 30-60 ns per lookup and 2.2MB of memory for the largest set of nodes.

Table 3 shows the peak-to-average load ratio for multi-probe consistent hash with $K = 21$. The peak-to-average load ratio converges to 1.05.

Table 4 shows the peak-to-average load ratio for ring consistent hash with $J = \ln n$ hashes per node. The peak-to-average load ratio converges to e .

Table 3: Peak-to-average, $K = 21$

Number of nodes	median of trials	90%ile of trials	99%ile of trials
10	1.04	1.13	1.24
100	1.05	1.08	1.10
1,000	1.05	1.06	1.07
10,000	1.05	1.06	1.06
100,000	1.05	1.06	1.06

Table 4: Peak-to-average, $J = \ln n$

Number of nodes	J	median of trials	90%ile of trials	99%ile of trials
10	2	2.23	3.05	3.96
100	4	2.64	3.24	4.05
1,000	6	2.84	3.29	3.75
10,000	9	2.79	3.11	3.51
100,000	11	2.89	3.15	3.40

Table 5: Peak-to-average, $J = 700 \ln n$

Number of nodes	J	median of trials	90%ile of trials	99%ile of trials
10	1611	1.04	1.06	1.08
100	3223	1.05	1.06	1.07
1,000	4835	1.05	1.05	1.06
10,000	6447	1.05	1.05	1.06
100,000	-	-	-	-

Table 6: Initialization time (ns) per node

Number of nodes	Multi-probe c. h. $\varepsilon = 0.05$	Ring c. h. $\varepsilon = 0.05$	Jump c. h. $\varepsilon = 0$
10	28	58,000	0
100	29	175,000	0
1,000	31	555,000	0
10,000	41	910,000	0
100,000	40	-	0

Table 5 shows the peak-to-average load ratio for ring consistent hash with $J = 700 \ln n$ hashes per node. The peak-to-average load ratio converges to 1.05. At 10,000 nodes the table required 1,400MB of memory. At 100,000 nodes the table did not fit in memory on the available machines.

5.2 Timings

For multi-probe consistent hash we set $K = 21$, obtaining a peak-to-average load ratio of 1.05.

For ring consistent hash we set $J = 700 \ln n$, obtaining a peak-to-average load ratio of 1.05. The implementation of ring consistent hash uses a hash table for $O(1)$ assignment, similar to the implementation of multi-probe consistent hash.

The implementation of jump consistent hash is taken without modification from [5]. Jump consistent hash achieves a peak-to-average load ratio of 1.0.

All implementations are in C++. Binaries are compiled on a 64-bit platform using GNU C++ and measured on 1 core of an Intel Xeon W3690 @3.47GHz.

Table 6 shows the initialization time per node. Multi-probe consistent hash is constant except for a step at 10,000 nodes, as the hash table spills to L3 cache². Ring consistent hash requires orders of magnitude more initialization time per node. Jump consistent hash requires no initialization.

Table 7 shows the memory per node, where we have used 64 bit hashes and 64 bit node identifiers. Multi-probe consistent hash uses constant memory per node. Ring consistent hash requires orders of magnitude more memory, commensurate with its high initialization time. Jump consistent hash requires no memory.

Table 8 shows the time per key. Multi-probe consistent hash is takes constant time modulo

²Cache spilling is visible throughout the timings. We will not comment upon each instance.

Table 7: Memory (bytes) per node

Number of nodes	Multi-probe c. h. $\varepsilon = 0.05$	Ring c. h. $\varepsilon = 0.05$	Jump c. h. $\varepsilon = 0$
10	22	35,000	0
100	22	71,000	0
1,000	22	106,000	0
10,000	22	142,000	0

Table 8: Assignment time (ns)

Number of nodes	Multi-probe c. h. $\varepsilon = 0.05$	Ring c. h. $\varepsilon = 0.05$	Jump c. h. $\varepsilon = 0$
10	350	29	32
100	420	60	50
1,000	430	110	67
10,000	590	130	80
100,000	590	-	94

cache effects. Ring consistent hash is a few times faster. Jump consistent hash is generally fastest as it does not access memory.

Table 9 shows the amortized time per insertion or removal of a node, measured by inserting from empty to full then removing from full to empty again (in random order). Multi-probe consistent hash requires only $O(1)$ amortized time per insertion or removal. Ring consistent hash requires orders of magnitude more time per update. Jump consistent hash requires no time to update, as it does not maintain a hash table.

It's important to note that all timings above are for uncontended caches, such that the hash table of nodes are cached near the CPU. However caches are typically contended in practice, which may evict the hash table of nodes to L3 or even main memory. Key assignment and node updates may be commensurately slower for multi-probe and ring

Table 9: Update time (ns)

Number of nodes	Multi-probe c. h. $\varepsilon = 0.05$	Ring c. h. $\varepsilon = 0.05$	Jump c. h. $\varepsilon = 0$
10	33	135,000	0
100	51	360,000	0
1000	70	1,000,000	0
10000	79	1,800,000	0
100000	107	-	0

consistent hash.

6 Discussion

Jump consistent hash is not generally applicable, as it cannot handle the loss of an arbitrary node. However where applicable it generally requires less time and space than the alternatives, in which case we recommend jump consistent hash.

Ring consistent hash has fast key assignment: just one hash table lookup. However to achieve a peak-to-average load ratio of $1 + \varepsilon$ over n nodes it requires $O(\frac{n \ln n}{\varepsilon^2})$ memory, potentially multiple gigabytes in practice. It is correspondingly slow to initialize and to update.

Multi-probe consistent hash stores each node just once in a hash table, so it requires only $O(n)$ memory and supports updates in $O(1)$ expected amortized time. To achieve a peak-to-average load ratio of $1 + \varepsilon$ it requires $O(\frac{1}{\varepsilon})$ time per lookup. In practice it can achieve a peak-to-average load ratio of 1.05 in 350-600 ns per key assignment, while scaling to larger node sets than possible with ring consistent hash. This makes multi-probe consistent hash an attractive replacement for ring consistent hash.

It's interesting to note the similarity between multi-probe consistent hash and cuckoo hashing [8], in which hashing keys two ways achieves a load factor up to $\frac{1}{2}$ for an in-memory hash table. The authors speculate that there might be fruitful connections to explore here.

7 Acknowledgements

The authors would like to thank Eric Veach, Raul Vera and David Symonds for many insightful comments on the draft paper.

References

- [1] H Cramèr. Sur un nouveau théorème-limite de la théorie des probabilités. In *Actualités Scientifiques et Industrielles*, volume 736, pages 5–23, 1938.
- [2] Giuseppe DeCandia, Deniz Hastorun, Madan Jampani, Gunavardhan Kakulapati, Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Voss, and Werner Vogels. Dynamo: Amazon's highly available key-value store. In *Proceedings of*

Twenty-first ACM SIGOPS Symposium on Operating Systems Principles, SOSP '07, pages 205–220, New York, NY, USA, 2007. ACM.

- [3] William Feller. *An introduction to probability theory and its applications. Vol. II.* Second edition. John Wiley & Sons Inc., New York, 1971.
- [4] David Karger, Eric Lehman, Tom Leighton, Rina Panigrahy, Matthew Levine, and Daniel Lewin. Consistent hashing and random trees: Distributed caching protocols for relieving hot spots on the world wide web. In *Proceedings of the Twenty-ninth Annual ACM Symposium on Theory of Computing*, STOC '97, pages 654–663, New York, NY, USA, 1997. ACM.
- [5] John Lamping and Eric Veach. A fast, minimal memory, consistent hash algorithm. *CoRR*, abs/1406.2294, 2014.
- [6] Petar Maymounkov and David Mazières. Kademlia: A peer-to-peer information system based on the xor metric. In *Revised Papers from the First International Workshop on Peer-to-Peer Systems*, IPTPS '01, pages 53–65, London, UK, UK, 2002. Springer-Verlag.
- [7] Colin McDiarmid. Concentration. In *Probabilistic Methods for Algorithmic Discrete Mathematics*, volume 16 of *Algorithms and Combinatorics*, pages 195–248. Springer Berlin Heidelberg, 1998.
- [8] Rasmus Pagh and Flemming Friche Rodler. Cuckoo hashing. *J. Algorithms*, 51(2):122–144, May 2004.
- [9] Christian Schindelhauer and Heinz Nixdorf. Weighted distributed hash tables. In *Proceedings of the 17th ACM Symposium on Parallel Algorithms and Architectures (SPAA)*, page 218, 2005.
- [10] David G. Thaler and China V. Ravishankar. Using name-based mappings to increase hit rates. *IEEE/ACM Transactions on Networking*, 6(1), February 1998.
- [11] Wei Wang and China V. Ravishankar. Hash-based virtual hierarchies for scalable location service in mobile ad-hoc networks. *Mob. Netw. Appl.*, 14(5):625–637, October 2009.

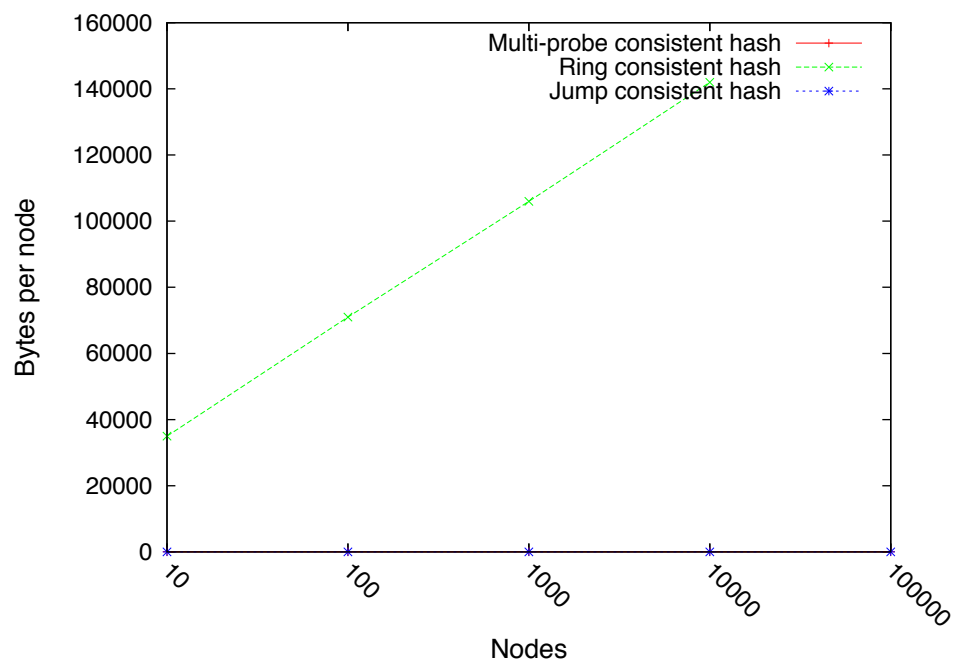


Figure 1: Memory (bytes) per node, $\varepsilon \leq 0.05$

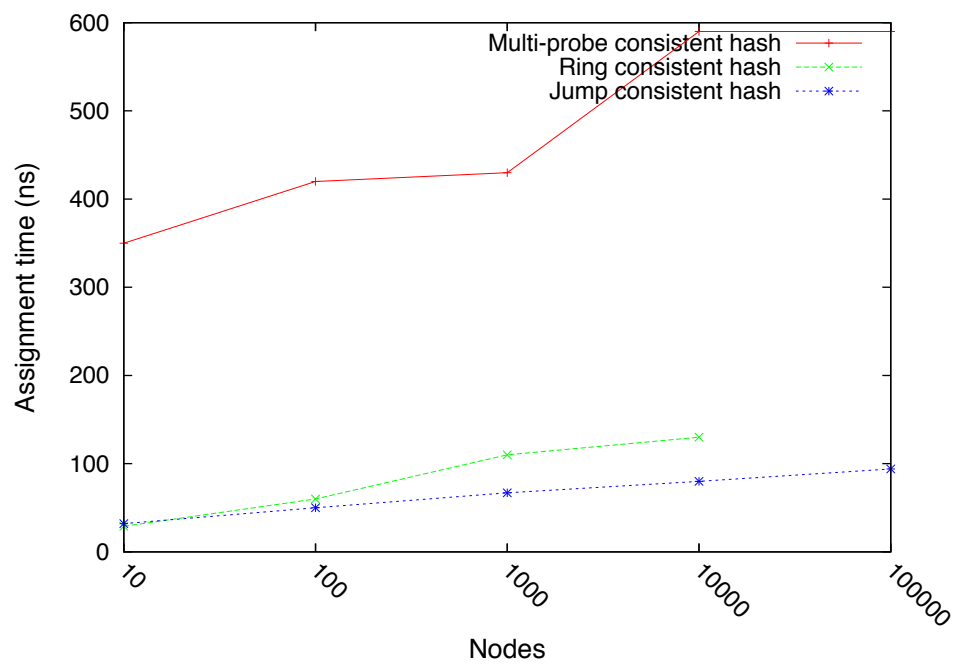


Figure 2: Assignment time (ns) per key, $\varepsilon \leq 0.05$

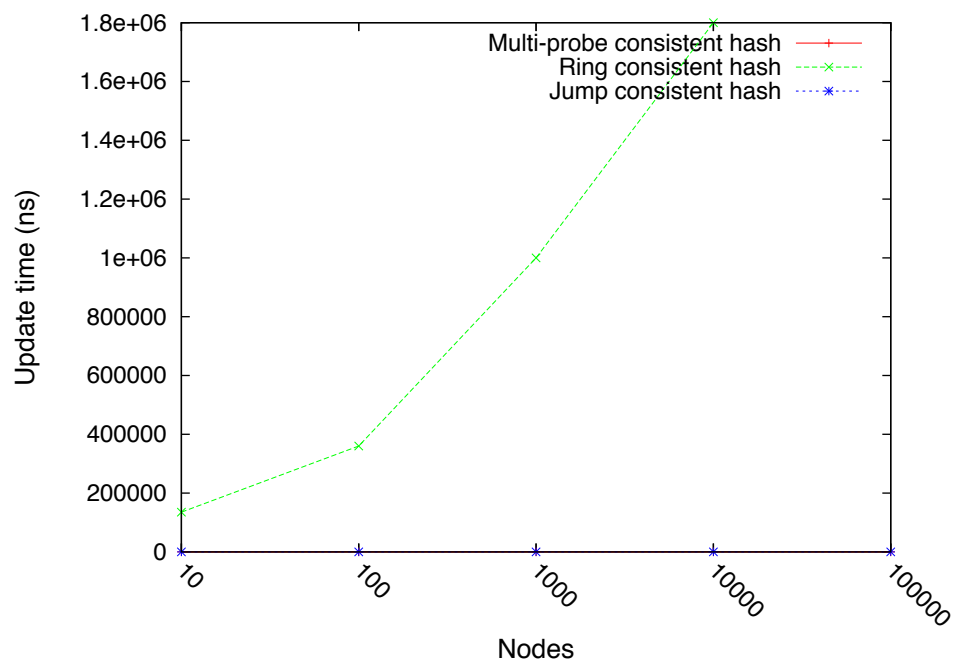


Figure 3: Update time (ns), $\varepsilon \leq 0.05$